

(Not yet Adaptive) Compression of In-Memory Databases

Database Implementation Lab Course

Leon Windheuser

March 14, 2023

Project Introduction

Goal: We want to compress the transient part of DuckDB

Project Introduction

Goal: We want to compress the transient part of DuckDB

- ▶ Open Source SQL OLAP RDBMS in-process developed in Amsterdam research center CWI (SQLite for OLAP)
<https://github.com/duckdb/duckdb>
- ▶ Columnar storage format
- ▶ Vectorized execution engine
- ▶ Has already lots of different compression possibilities for persistent data on disk

Project Introduction

Goal: We want to compress the transient part of DuckDB

- ▶ Open Source SQL OLAP RDBMS in-process developed in Amsterdam research center CWI (SQLite for OLAP)
<https://github.com/duckdb/duckdb>
- ▶ Columnar storage format
- ▶ Vectorized execution engine
- ▶ Has already lots of different compression possibilities for persistent data on disk

How do we compress the transient data while having efficient lookups without decompressing everything?

Background: Succinct Data Structures


- ▶ Data structures that use space close to the theoretic lower bound but allow efficient query operations (in place without needing to decompress)
- ▶ Exists e.g. (bit) vectors, trees, planar graphs, ...

Background: Succinct Data Structures

- ▶ Data structures that use space close to the theoretic lower bound but allow efficient query operations (in place without needing to decompress)
- ▶ Exists e.g. **(bit) vectors**, trees, planar graphs, ...


Succinct Integer Vector

Space requirement for integer x is $\ell = \lceil \log_2(x) \rceil$ bits




Succinct Integer Vector

Space requirement for integer x is $\ell = \lceil \log_2(x) \rceil$ bits



Encode integers with the minimal length of the max integer
 $3 = \lceil \log_2(7) \rceil$



We already reduce memory by 25%

SDSL: Succinct Data Structure Library

- ▶ C++11 library for succinct data structures
- ▶ Open Source <https://github.com/simongog/sdsl-lite>
- ▶ Contains a variety of succinct data structures. We use **Integer Vector** to compress integer columns.

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

Reduces memory by 56.2% (\approx 22.5 KB)

Delta Compression of sds1::int_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 24, size: 30008

Delta Compression of sds1::int_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

Delta Compression of sds1::int_vector

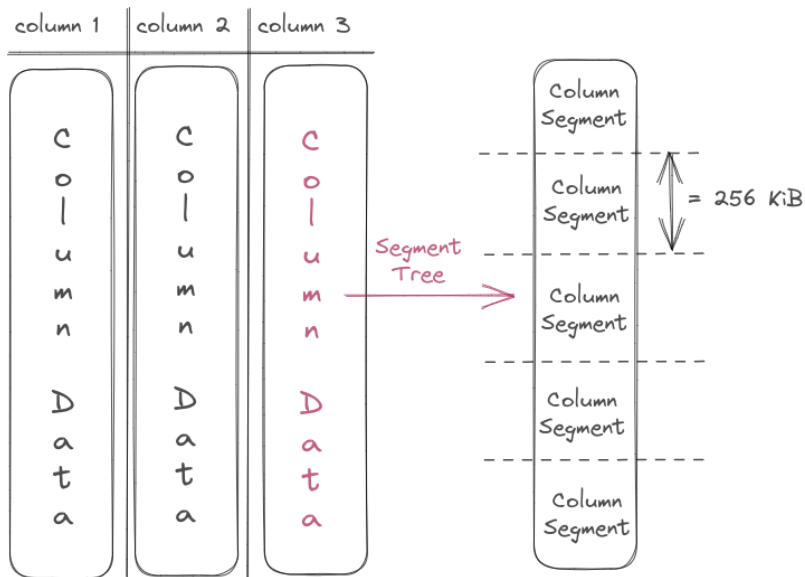
```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

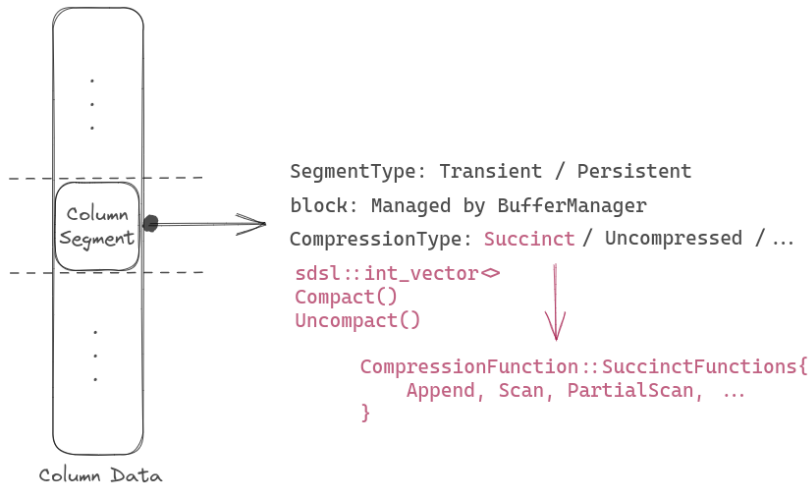
Width: 14, size: 17513

In DuckDB, we know the minimum of the vector directly without searching (column statistics)

DuckDB Storage Architecture Bird's Eye View (100 meters)




DuckDB Storage Architecture Bird's Eye View (10 meters)



Evaluation: Sequential Insert and Total Scan

Scanning 1M rows.


```
SELECT * FROM t1;
```



Evaluation: Sequential Insert and Total Scan

Scanning 1M rows.


```
SELECT * FROM t1;
```



Evaluation: Zipf Selection

10000 selections with Zipf distribution of 1M total rows.


```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



Evaluation: Zipf Selection


10000 selections with Zipf distribution of 1M total rows.

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```




Evaluation: Zipf Out-Of-Memory (Limit 1GB)

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```

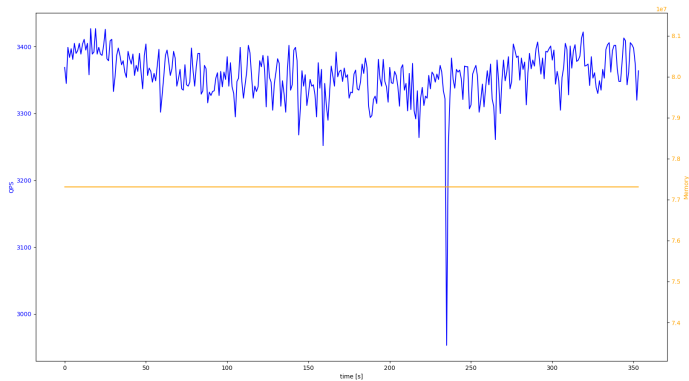


Evaluation: Zipf Out-Of-Memory (Limit 1GB)


```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```




Zipf Over Time: Not Succinct



Zipf Over Time: Succinct



Zipf Over Time: Adaptive Succinct



Conclusion

- ▶ For OLAPish queries: It is not (yet?) worth it (large overhead due to copying).
- ▶ For OLTP transactions: It might be worth it. Reduces memory by $\approx 40\%$ but increases runtime by $\approx 35\%$.
- ▶ Great benefit if succinct representation fits in memory vs. spilling to disk.

Future Work and Discussion

1. Copying and shifting data is most time consuming ($\approx 40\%$) since execution engine expects a flat “normal” vector.
 - ▶ Non-succinct passes its data pointer that we need to decompress and copy the data.
 - ▶ Unnecessary since we still support random access and operations needed for the execution engine.
 - ▶ Non-succinct data pointer used everywhere in the execution engine (> 300 appearances). **Rewrite necessary?**
2. Adaptive compression for rarely accessed segments. Zipf Distribution accesses 4/50 segments over 70% of the time.
 - ▶ How to track access statistics over time for segments?
 - ▶ What if the access statistics change after a greater period of time?