

(Not yet Adaptive) Compression of In-Memory Databases

Database Implementation Lab Course

Leon Windheuser

February 28, 2023

Project Introduction

Compression of the In-Memory part of DuckDB

Project Introduction

Compression of the In-Memory part of DuckDB

- ▶ Open Source SQL OLAP RDBMS developed in Amsterdam research centre CWI
(<https://github.com/duckdb/duckdb>)
- ▶ Columnar Storage format
- ▶ Vectorized execution engine
- ▶ Has lots of different compression possibilities for persistent data on disk

Project Introduction

Compression of the In-Memory part of DuckDB

- ▶ Open Source SQL OLAP RDBMS developed in Amsterdam research centre CWI
(<https://github.com/duckdb/duckdb>)
- ▶ Columnar Storage format
- ▶ Vectorized execution engine
- ▶ Has lots of different compression possibilities for persistent data on disk

How do we compress data in-memory while having efficient lookups without decompressing everything?

Succinct Data Structures

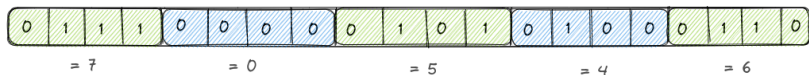
- ▶ Data structures which uses close to the *information-theoretic* lower bound of space but allows efficient query operations (in-place without needing to decompress)
- ▶ Exists for e.g. (bit) vectors, trees, planar graphs, ...

Succinct Data Structures

- ▶ Data structures which uses close to the *information-theoretic* lower bound of space but allows efficient query operations (in-place without needing to decompress)
- ▶ Exists for e.g. **(bit) vectors**, trees, planar graphs, ...


Succinct Integer Vector

Space requirement for integer x is $\ell = \lfloor \log_2(x) \rfloor + 1$ bits




Succinct Integer Vector

Space requirement for integer x is $\ell = \lfloor \log_2(x) \rfloor + 1$ bits



Encode integers with the minimal length of the max integer

$$3 = \lfloor \log_2(7) \rfloor + 1$$



We already reduce memory by 25%

SDSL: Succinct Data Structure Library

- ▶ C++11 library and abstraction for succinct data structures
- ▶ Open Source <https://github.com/simongog/sdsl-lite>
- ▶ Contains variety of different data structures. For now we only used the **Integer Vectors**.

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

Reduces memory by 56.2% (\approx 22.5 KB)

Extract min element from sds1::int_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 24, size: 30008

Extract min element from sds1::int_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

Extract min element from sds1::int_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

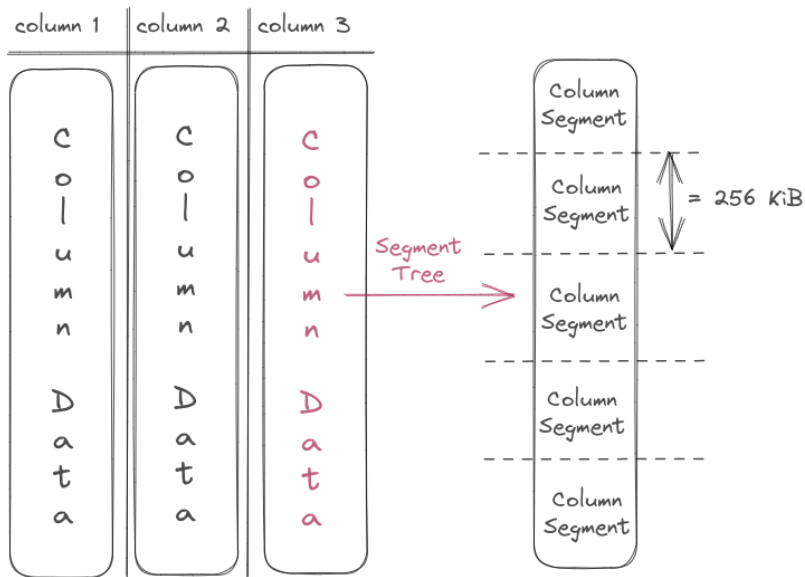
Width: 14, size: 17513

In DuckDB we know the minimum of the vector directly without searching (column statistics)

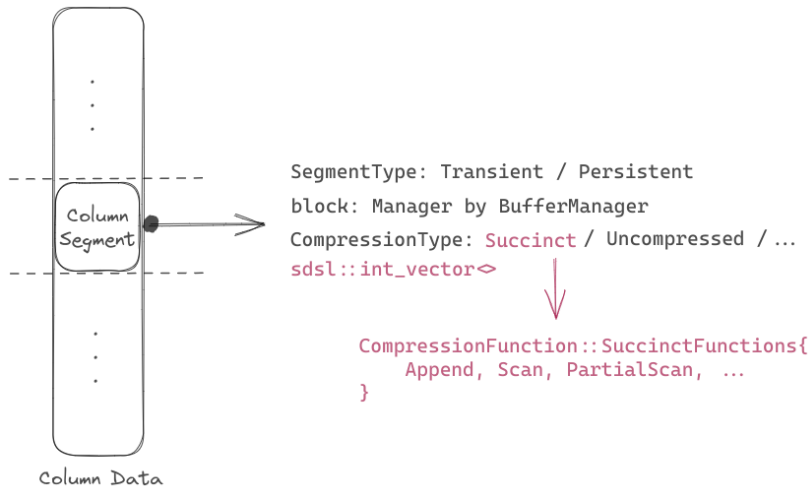
Random Access of `sds1::int_vector[i]` vs
`std::vector[i]`

Byte-Align `sdsl::int_vector`

DuckDB Storage Architecture 100 meter view




DuckDB Storage Architecture 10 meter view



Benchmarks: Sequential Insert and total Scan

Scanning 10^6 rows.


```
SELECT * FROM t1;
```



Benchmarks: Sequential Insert and total Scan

10.000 selections with Zipf Distribution of 10^6 total rows.


```
SELECT * FROM t1;
```



Benchmarks: Zipf Selection

10.000 selections with Zipf Distribution of 10^6 total rows.

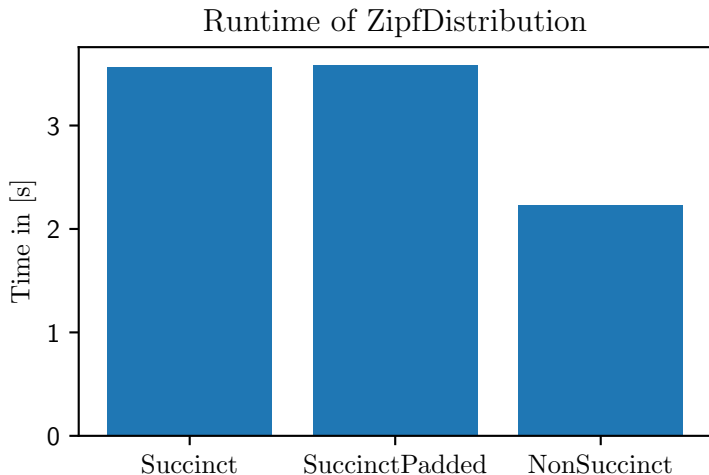
```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



Benchmarks: Zipf Selection


10.000 selections with Zipf Distribution of 10^6 total rows.

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```




Benchmarks: Zipf Out Of Memory (Limit 1GB)

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



Benchmarks: Zipf Out Of Memory (Limit 1GB)

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



Conclusion and Future Work

- ▶ For OLAPish queries it is not (yet?) worth it, large overhead.
- ▶ For OLTP transactions it might be worth it. Reduces memory by $\approx 40\%$ but increases runtime by $\approx 35\%$.
- ▶ Huge benefit if succinct representation fits in memory vs spilling to disk.

Conclusion and Future Work

- ▶ For OLAPish queries it is not (yet?) worth it, large overhead.
- ▶ For OLTP transactions it might be worth it. Reduces memory by $\approx 40\%$ but increases runtime by $\approx 35\%$.
- ▶ Huge benefit if succinct representation fits in memory vs spilling to disk.

TODO:

- ▶ Need to investigate in access statistics to only compress rarely used column segments.
- ▶ Currently copying a lot of data as execution engine expects an uncompressed flat vector. Adapt execution engine to allow succinct vectors as well since they allow random access.