

# (Not yet Adaptive) Compression of In-Memory Databases

Database Implementation Lab Course

Leon Windheuser

March 14, 2023

# Project Introduction

We want to compress the transient part of DuckDB

# Project Introduction

We want to compress the transient part of DuckDB

- ▶ Open Source SQL OLAP RDBMS in-process developed in Amsterdam research centre CWI (SQLite for OLAP)  
<https://github.com/duckdb/duckdb>
- ▶ Columnar Storage format
- ▶ Vectorized execution engine
- ▶ Has already lots of different compression possibilities for persistent data on disk

# Project Introduction

We want to compress the transient part of DuckDB

- ▶ Open Source SQL OLAP RDBMS in-process developed in Amsterdam research centre CWI (SQLite for OLAP)  
<https://github.com/duckdb/duckdb>
- ▶ Columnar Storage format
- ▶ Vectorized execution engine
- ▶ Has already lots of different compression possibilities for persistent data on disk

**How do we compress the transient data while having efficient lookups without decompressing everything?**

# Background: Succinct Data Structures

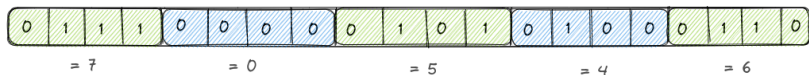
- ▶ Data structures that use space close to the theoretic lower bound but allow efficient query operations (in place without needing to decompress)
- ▶ Exists e.g. (bit) vectors, trees, planar graphs, ...

# Background: Succinct Data Structures

- ▶ Data structures that use space close to the theoretic lower bound but allow efficient query operations (in place without needing to decompress)
- ▶ Exists e.g. **(bit) vectors**, trees, planar graphs, ...


# Succinct Integer Vector

Space requirement for integer  $x$  is  $\ell = \lceil \log_2(x) \rceil$  bits




# Succinct Integer Vector

Space requirement for integer  $x$  is  $\ell = \lceil \log_2(x) \rceil$  bits



Encode integers with the minimal length of the max integer  
 $3 = \lceil \log_2(7) \rceil$



**We already reduce memory by 25%**



# SDSL: Succinct Data Structure Library

- ▶ C++11 library and abstraction for succinct data structures
- ▶ Open Source <https://github.com/simongog/sdsl-lite>
- ▶ Contains variety of succinct data structures. We use their **Integer Vector** interface to compress integers.

## SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

## SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

## SDSL: Integer Vectors

```
sdsl::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++) v[i] = i;  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;  
sdsl::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sdsl::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

Reduces memory by 56.2% ( $\approx$  22.5 KB)

## Delta Compression of sds1::int\_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
      << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 24, size: 30008

## Delta Compression of sds1::int\_vector

```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

## Delta Compression of sds1::int\_vector

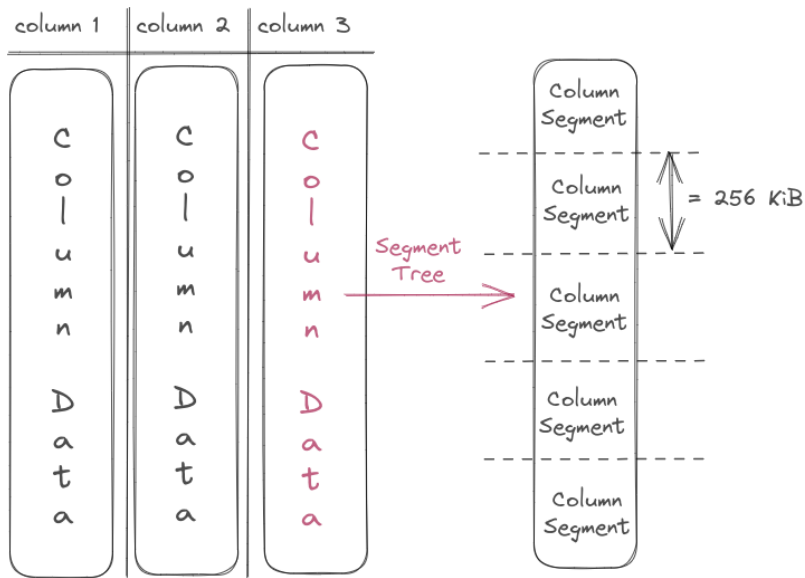
```
sds1::int_vector<32> v(10000);  
for (size_t i = 0; i < 10000; i++)  
    v[i] = i + 10.000.000;  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;  
extractMinFromVector(v);  
sds1::util::bit_compress(v);  
cout << "Width: " << v.width() << ", size: "  
    << sds1::size_in_bytes(v) << endl;
```

Width: 32, size: 40008

Width: 14, size: 17513

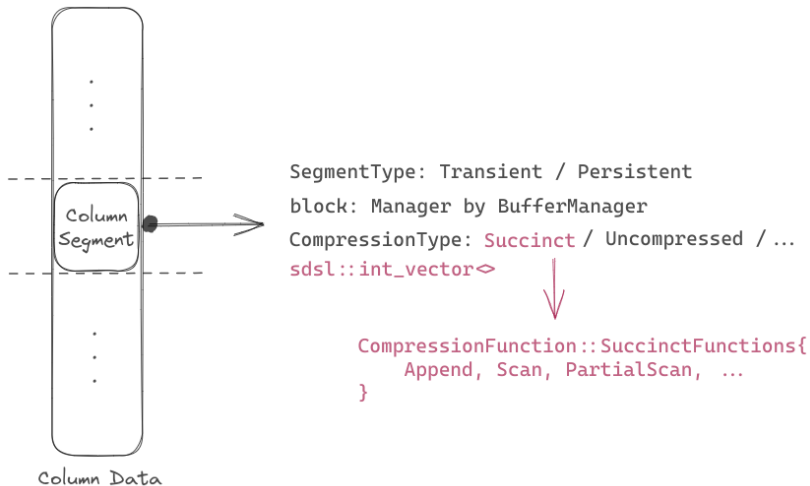
In DuckDB we know the minimum of the vector directly without searching (column statistics)

# DuckDB Storage Architecture Bird's Eye View (100 meter)






# DuckDB Storage Architecture Bird's Eye View (10 meter)



## Evaluation: Sequential Insert and Total Scan

Scanning  $10^6$  rows.

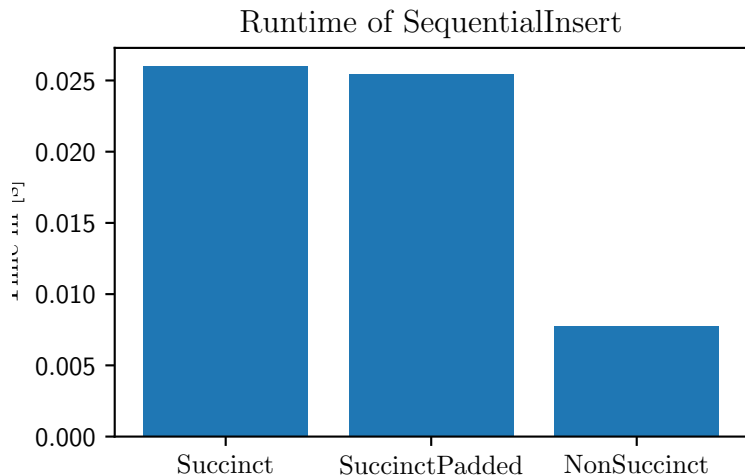
```
SELECT * FROM t1;
```



# Evaluation: Sequential Insert and Total Scan

Scanning  $10^6$  rows.


```
SELECT * FROM t1;
```



## Evaluation: Zipf Selection

10.000 selections with Zipf Distribution of  $10^6$  total rows.

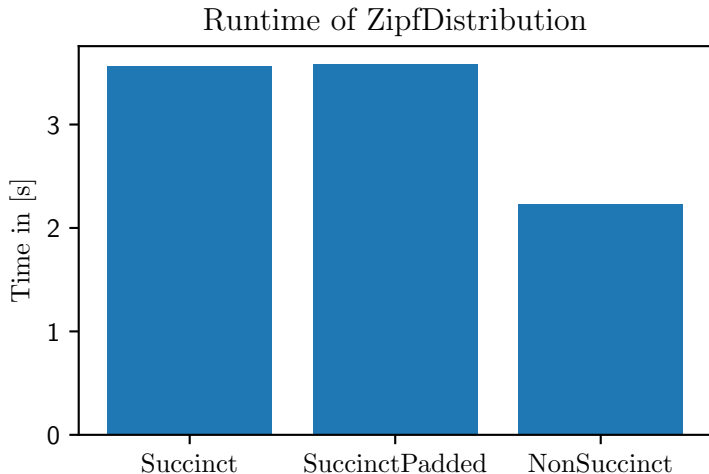
```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



## Evaluation: Zipf Selection

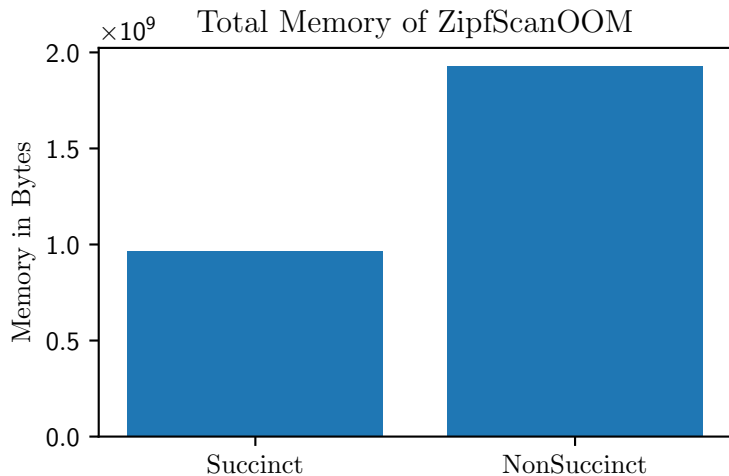
10.000 selections with Zipf Distribution of  $10^6$  total rows.

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```




## Evaluation: Zipf Out-Of-Memory (Limit 1GB)

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



## Benchmarks: Zipf Out-Of-Memory (Limit 1GB)

```
SELECT i FROM t1  
WHERE i == {ZIPF_DISTRIBUTED_NUMBER};
```



# Conclusion

- ▶ For OLAPish queries it is not (yet?) worth it, large overhead.
- ▶ For OLTP transactions it might be worth it. Reduces memory by  $\approx 40\%$  but increases runtime by  $\approx 35\%$ .
- ▶ Huge benefit if succinct representation fits in memory vs spilling to disk.



# Future Work and Discussion

1. Copying and shifting data is most time consuming ( $\approx 40\%$ ) since execution engine expects a flat "normal" vector.
  - ▶ Non succinct passes its data pointer, we need to decompress and copy the data.
  - ▶ Unecessary, since we still support random access and operations needed for the execution engine.
  - ▶ Non succinct data pointer used everywhere in the execution engine ( $> 300$  appearances). **Rewrite necessary?**
2. Adaptive compression for rarely accessed segments. Zipf Distribution accesses 4/50 segments over 70% of the time.
  - ▶ How to track access statistics over time for segments?
  - ▶ What if the access statistics change after greater period of time?